



## Mercari 1st place solution

Konstantin Lopuhin & Pawel Jankiewicz

2018-03-31

# About us



**Konstantin Lopuhin**

Software engineer at Scrapinghub

Москва, Россия

Joined 6 years ago · last seen in the past day



## Competitions Master



Current Rank

**32**

of 80,984

Highest Rank

**30**



4



1



0

[Mercari Price Suggestion C...](#)

🥇 · a month ago · Top 1%

1<sup>st</sup>

of 2384

[NOAA Fisheries Steller Sea...](#)

🥇 · 9 months ago · Top 1%

2<sup>nd</sup>

of 385

[Dstl Satellite Imagery Feat...](#)

🥇 · a year ago · Top 2%

5<sup>th</sup>

of 419



**Paweł Jankiewicz**

Warszawa, mazowieckie, Polska

Joined 6 years ago · last seen in the past day

[in](#) <http://logicalai.io>

## Competitions Grandmaster



Current Rank

**76**

of 80,984

Highest Rank

**40**



9



2



3

[Mercari Price Suggestion C...](#)

🥇 · a month ago · Top 1%

1<sup>st</sup>

of 2384

[Will I Stay or Will I Go?](#)

🥇 · 5 years ago · Top 9%

1<sup>st</sup>

of 12

[GE Flight Quest](#)

🥇 · 5 years ago · Top 2%

2<sup>nd</sup>

of 173

# We won!

At the time of merging we were 1st (Konstantin) and 2nd (Pawel).

#	$\Delta$ pub	Team Name	Kernel	Team Members	Score
1	—	Pawel and Konstantin		 ★★★★★	0.37758
2	—	Mercaring (Nima & Chahhou)		 ★★★★★	0.38875
3	▲1	bird		RUA ★★★★★	0.39134
4	▲1	Chenglong Chen		 ★★★★★	0.39299
5	▲2	anttip		 ★★★★★	0.39603
6	▲5	Fair trade		 ★★★★★	0.39713
7	▲3	Basil		 ★★★★★	0.39720
8	—	RDizzl3 and Sergei		 ★★★★★	0.39734
9	▲3	LeeYun	<> ensemble model	 ★★★★★	0.39752
10	▲3	stooging the stooges		 ★★★★★	0.39766

# Mercari competition

<https://www.kaggle.com/c/mercari-price-suggestion-challenge/>



**NWT JUICY COUTURE JACKET XS**

California • 03/23/2018 06:45 PM • [Report item](#)

**\$ 45.00**  0

[Sign up now and buy at \\$ 40.00](#)

[Buy now](#)

Condition  Like new

Size  XS (0-2)

Shipping  \$7.00

Brand Juicy Couture

Category [Women](#)  
[Athletic apparel](#)  
[Jackets](#)

Description

Brand new!! JUICY COUTURE JACKET XS!!!



# Evaluation RMSLE

## RMSLE


$$\epsilon = \sqrt{\frac{1}{n} \sum_{i=1}^n (\log(p_i + 1) - \log(a_i + 1))^2}$$

Better to convert to RMSE - optimize directly


```
y_log = log(y + 1)
model.fit(X, y_log)
prediction_log = model.predict(X_test)
prediction = exp(prediction_log) - 1
```

# Why Logarithm?

Price distribution



Log(price) distribution



# Code competition

Only 60 minutes to train and predict!

System specification: 16 GB ram, 1 GB disk,  $\sim 4$  cores

## Advantages

- ▶ No huge ensembles
- ▶ Team collaboration - common code base
- ▶ Small models, fast iteration

## Disadvantages

- ▶ Pure optimization
- ▶ Unstable platform
- ▶ 2 stage competition (5x bigger test data)

# Build system

<< 8f5b46e

Draft saved Python

Commit & Run

Code	Data	Settings	Versions
------	------	----------	----------

```
1 import gzip
2 import base64
3 import os
4 import sys
5 sys.path.append('/kaggle/working')
6
7 # this is base64 encoded source code
8 file_data = {'__init__.py': 'H4sIAMxEgVoC/wMAAAAAAAAAA=', 'config.py': 'H4sIAMxEgVoC/4VUXW+bMBR996+w0iFAahht162rLEk0eA1aQq0QTnteLBRuqFXAFXbabtP++y4QMmJS
9
10 os.system('mkdir mercari')
11
12
13 for fn, encoded in file_data.items():
14     print(fn)
15     with open('mercari/' + fn, 'wb') as out:
16         out.write(gzip.decompress(base64.b64decode(encoded)))
17
18
19 with open('setup.py', 'wt') as out:
20     .
    .
    .
```

▶

Console

Environment Variables

CPU 0% RAM 215MB/17.2GB Disk 55.4MB/1GB

[Initializing]  
[Starting]  
[Running]  
Your kernel is now running in the cloud. Here are some things you can do with it:  
\* Use the Play button or [SHIFT]+[ENTER] to execute the current line of your script (or whatever's highlighted).  
\* Enter some code at the bottom of this Console tab and press [ENTER].



## About the data

1.5 million observations in the training data

Column name	Type	# unique values
name	text	-
item_description	text	-
item_condition_id	categorical - ordinal	5
category_name	text/categorical	1288
brand_name	text/categorical	4810
shipping	boolean	-

## Example Item

```
{  
  "name": "NYX GLITTER GLUE+GLITTER BUNDLE",  
  "item_description": "FREE FAST SHIPPING EXTRA FREE☆☆ BEAUTY GIFTS:)  
;;♡♡☆☆☆ Extra free skincare gift:));;♡♡☆☆◇◇◇ NYX GLITTER PRIMER glue..  
brand new sealed in box.. full size 4 NYX eye. glitters!!brand new..all in 5  
grams jars each.. variety of colors to match just about any makeup look you  
decide to create Wet n wild small concealer brush.. brand new and sealed..  
perfect sized brush to apply glitter primer plus glitters onto your eyelid  
without being messy.. works a lot like the real techniques detailer brush",  
  "item_condition_id": 1,  
  "category_name": "Beauty/Makeup/Eyes",  
  "brand_name": "NYX",  
  "shipping": 1,  
  "price": 28.0  
}
```

# Data preprocessing: Declarative vs Imperative

## Imperative

```
D vect = CountVectorizer()  
A vect.fit(X)  
A mat = vect.transform(X)  
D rf = RandomForestRegressor()  
A rf.fit(mat, y)
```

D = declaration, A = action

## Declarative

```
D model = make_pipeline(  
D     CountVectorizer(),  
D     RandomForestRegressor()  
D )  
A model.fit(X, y)
```

## "It's pipelines all the way down"

```
15 def prepare_vectorizer_1_tf(n_jobs=4):
16     tokenizer = FastTokenizer()
17     vectorizer = make_pipeline(
18         FillEmpty(),
19         PreprocessDataPJ(n_jobs=n_jobs),
20         make_union_mp(
21             make_pipeline(
22                 PandasSelector(columns=['name', 'item_description']),
23                 ConcatTexts(columns=['name', 'item_description'],
24                                 use_separators=True),
25                 PandasSelector(columns=['text_concat']),
26                 CountVectorizer(ngram_range=(1, 1), binary=True, min_df=5, tokenizer=tokenizer, dtype=np.float32)
27             ),
28             make_pipeline(PandasSelector(columns=['category_name_clean']),
29                           CountVectorizer(tokenizer=tokenizer,
30                                           binary=True,
31                                           min_df=5,
32                                           dtype=np.float32)),
33             make_pipeline(PandasSelector(columns=['shipping', 'item_condition_id', 'brand_name_clean',
34                                                 'cat_1', 'cat_2', 'cat_3', 'no_cat']),
35                           PandasToRecords(),
36                           DictVectorizer(dtype=np.float32)),
37             n_jobs=n_jobs
38         ),
39         SparseMatrixOptimize(),
40         SanitizeSparseMatrix(),
41         ReportShape()
42     )
43     return vectorizer
```

# Preprocessing



- ▶ Text preprocessing - stemming
- ▶ Bag of words - 1,2-grams (with/without Tf-Idf)
- ▶ One hot encoding for categorical columns

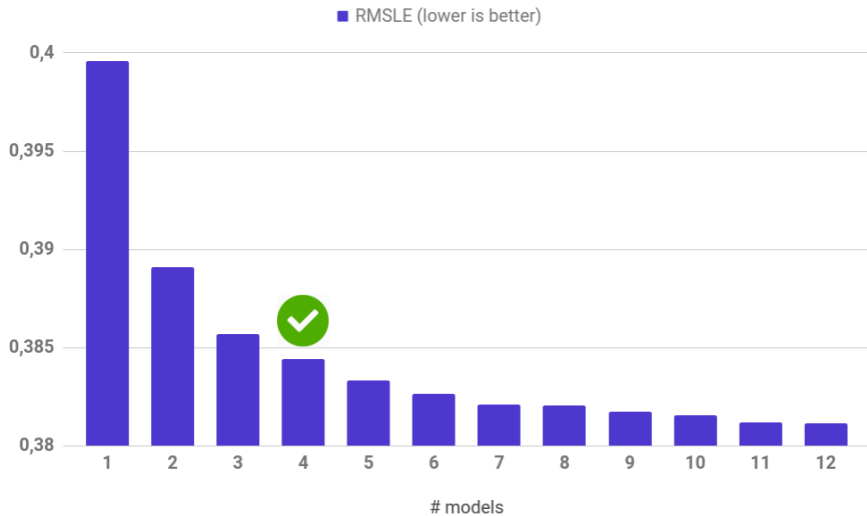


- ▶ Bag of character 3-grams

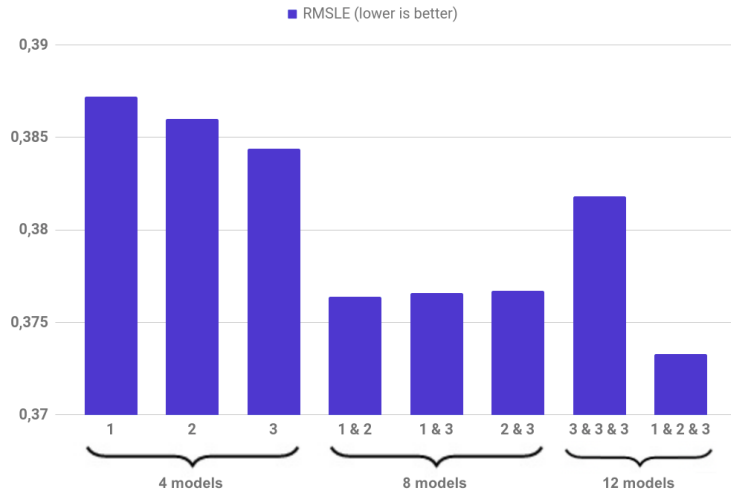


- ▶ Joining name, brand name and description into a single field
- ▶ NumericalVectorizer - vectorizing words using preceding numbers


# Why ensemble?



# Why 3 datasets?



## Our progress






# To Deep Learn or not to Deep Learn?

	Learning	Deep Learning
Architecture	<b>MLP</b>	LSTM, CNN
Activation	Tanh	<b>ReLU</b>
Optimization	SGD	<b>ADAM</b>
Categorical features	<b>One-hot encoding</b>	Embeddings

Workhorse model: sparse MLP (feedforward neural network)




# Why MLP?

- ▶ Fast to train: can afford hidden size 256 instead of 32–64 for RNN or Conv1D.
- ▶ Captures interactions between text and categorical features.
- ▶ Huge variance gives a strong ensemble with a single model type.


# Training

Adam, double batch size after each epoch, overfit



# Training


Adam, double batch size after each epoch, overfit, profit!





# Tricks

- ▶ Huber loss
- ▶ Regression via. classification
- ▶ Cheap feature binarization

# Huber Loss



# Regression via Classification





# Cheap feature binarization


TF-IDF features  $\Rightarrow$  Binary features

```
def feed_dict(self, X, binary_X=False):  
    coo = X.tocoo()  
    return {  
        self.indices: np.stack([coo.row, coo.col]).T,  
        self.values: coo.data if not binary_X else np.ones_like(coo.data),  
        self.shape: np.array(X.shape),  
    }
```

# Sparse MLP Implementation

- ▶ TensorFlow: `tf.sparse_tensor_dense_matmul`
- ▶ MXNet: `RowSparseNDArray`, sparse updates!
- ▶ Keras: `keras.Input(sparse=True)`
- ▶ Any framework: via embedding

## Optimization: One Model per Core



## Optimization: Memory

- ▶ TensorFlow: `threading`, `use_per_session_threads`
- ▶ MXNet: multiprocessing, memory efficient data loader

# Ensembling via Lasso

5% local validation, 1% on Kaggle. Very good LB correlation.

```
merge_predictions =  
-0.0203  
+0.0604 * data1_huber  
+0.1051 * data1_huber  
+0.0911 * data1_clf  
+0.0760 * data1_clf  
+0.0851 * data2_huber_bin  
+0.0981 * data2_huber  
+0.0819 * data2_clf_bin  
+0.0717 * data2_clf  
+0.0958 * data3_huber_bin  
+0.1226 * data3_huber  
+0.0578 * data3_clf_bin  
+0.0642 * data3_clf  
⇒ RMSLE 0.3733
```

# Didn't Work

- ▶ Grid Search
- ▶ Skip Connections
- ▶ Mixture of Experts
- ▶ Factorization Machines
- ▶ Fitting residuals

Code Golf: 0.3875 CV in 75 LOC, 1900 s

- ▶ Sparse MLP in Keras
- ▶ Train 4 models on 4 cores
- ▶ Custom preprocessing

```
1 import os; os.environ['OMP_NUM_THREADS'] = '1'
2 from contextlib import contextmanager
3 from functools import partial
4 from operator import itemgetter
5 from multiprocessing.pool import ThreadPool
6 import time
7 from typing import List, Dict
8
9 import keras as ks
10 import pandas as pd
11 import numpy as np
12 import tensorflow as tf
13 from sklearn.feature_extraction import DictVectorizer
14 from sklearn.feature_extraction.text import TfidfVectorizer as Tfidf
15 from sklearn.pipeline import make_pipeline, make_union, Pipeline
16 from sklearn.preprocessing import FunctionTransformer, StandardScaler
17 from sklearn.metrics import mean_squared_log_error
18 from sklearn.model_selection import KFold
19
20 @contextmanager
21 def timer(name):
22     t0 = time.time()
23     yield
24     print(f'[{name}] done in {time.time() - t0:.0f} s')
```

big CPU win!

boring  
stuff



```
26 def preprocess(df: pd.DataFrame) -> pd.DataFrame:
27     df['name'] = df['name'].fillna('') + ' ' + df['brand_name'].fillna('')
28     df['text'] = (df['item_description'].fillna('') + ' ' + df['name'] + ' ' +
29 df['category_name'].fillna(''))
30     return df[['name', 'text', 'shipping', 'item_condition_id']]
31
32 def on_field(f: str, *vec) -> Pipeline:
33     return make_pipeline(FunctionTransformer(itemgetter(f), validate=False), *vec)
34
```

```
35 def to_records(df: pd.DataFrame) -> List[Dict]:
36     return df.to_dict(orient='records')
```

```
37
38 def fit_predict(xs, y_train) -> np.ndarray:
39     X_train, X_test = xs
40     config = tf.ConfigProto(
41         intra_op_parallelism_threads=1, use_per_session_threads=1, inter_op_parallelism_threads=1)
42     with tf.Session(graph=tf.Graph(), config=config) as sess, timer('fit_predict'):
43         ks.backend.set_session(sess)
44         model_in = ks.Input(shape=(X_train.shape[1],), dtype='float32', sparse=True)
45         out = ks.layers.Dense(192, activation='relu')(model_in)
46         out = ks.layers.Dense(64, activation='relu')(out)
47         out = ks.layers.Dense(64, activation='relu')(out)
48         out = ks.layers.Dense(1)(out)
49         model = ks.Model(model_in, out)
50         model.compile(loss='mean_squared_error', optimizer=ks.optimizers.Adam(lr=3e-3))
51         for i in range(3):
52             with timer(f'epoch {i + 1}'):
53                 model.fit(x=X_train, y=y_train, batch_size=2**(11 + i), epochs=1, verbose=0)
54     return model.predict(X_test)[:, 0]
```

feature  
engineering

TF trick

← the MODEL

```

56 def main():
57     vectorizer = make_union(
58         on_field('name', Tfidf(max_features=100000, token_pattern='\\w+')),
59         on_field('text', Tfidf(max_features=100000, token_pattern='\\w+', ngram_range=(1, 2))),
60         on_field(['shipping', 'item_condition_id'],
61                 FunctionTransformer(to_records, validate=False), DictVectorizer()),
62         n_jobs=4)
63     y_scaler = StandardScaler()
64     with timer('process train'):
65         train = pd.read_table('../input/train.tsv')
66         train = train[train['price'] > 0].reset_index(drop=True)
67         cv = KFold(n_splits=20, shuffle=True, random_state=42)
68         train_ids, valid_ids = next(cv.split(train))
69         train, valid = train.iloc[train_ids], train.iloc[valid_ids]
70         y_train = y_scaler.fit_transform(np.log1p(train['price'].values.reshape(-1, 1)))
71         X_train = vectorizer.fit_transform(preprocess(train)).astype(np.float32)
72         print(f'X_train: {X_train.shape} of {X_train.dtype}')
73         del train
74     with timer('process valid'):
75         X_valid = vectorizer.transform(preprocess(valid)).astype(np.float32)
76     with ThreadPool(processes=4) as pool:
77         Xb_train, Xb_valid = [x.astype(np.bool).astype(np.float32) for x in [X_train, X_valid]]
78         xs = [[Xb_train, Xb_valid], [X_train, X_valid]] * 2
79         y_pred = np.mean(pool.map(partial(fit_predict, y_train=y_train), xs), axis=0)
80         y_pred = np.expm1(y_scaler.inverse_transform(y_pred.reshape(-1, 1))[:, 0])
81         print('Valid RMSLE: {:.4f}'.format(np.sqrt(mean_squared_log_error(valid['price'], y_pred))))

```

feature engineering

scale target

extra dataset  
for free

4X!

# Feature Engineering

```
26 def preprocess(df: pd.DataFrame) -> pd.DataFrame:
27     df['name'] = df['name'].fillna('') + ' ' + df['brand_name'].fillna('')
28     df['text'] = (df['item_description'].fillna('') + ' ' + df['name'] + ' ' +
29 df['category_name'].fillna(''))
30     return df[['name', 'text', 'shipping', 'item_condition_id']]
31
32     vectorizer = make_union(
33         on_field('name', Tfidf(max_features=100000, token_pattern='\w+')),
34         on_field('text', Tfidf(max_features=100000, token_pattern='\w+', ngram_range=(1, 2))),
35         on_field(['shipping', 'item_condition_id'],
36                 FunctionTransformer(to_records, validate=False), DictVectorizer()),
37         n_jobs=4)
38     y_scaler = StandardScaler()
```

join fields

more features!

# The Model

```
39 X_train, X_test = xs
40 config = tf.ConfigProto(
41     intra_op_parallelism_threads=1, use_per_session_threads=1, inter_op_parallelism_threads=1)
42 with tf.Session(graph=tf.Graph(), config=config) as sess, timer('fit_predict'):
43     ks.backend.set_session(sess)
44     model_in = ks.Input(shape=(X_train.shape[1],), dtype='float32', sparse=True)
45     out = ks.layers.Dense(192, activation='relu')(model_in)
46     out = ks.layers.Dense(64, activation='relu')(out)
47     out = ks.layers.Dense(64, activation='relu')(out)
48     out = ks.layers.Dense(1)(out)
49     model = ks.Model(model_in, out)
50     model.compile(loss='mean_squared_error', optimizer=ks.optimizers.Adam(lr=3e-3))
51     for i in range(3):
52         with timer(f'epoch {i + 1}'):
53             model.fit(x=X_train, y=y_train, batch_size=2**(11 + i), epochs=1, verbose=0)
54     return model.predict(X_test)[: , 0]
```

TF trick →

MLP

increase batch size

# Other Solutions

Popular models:

- ▶ Ridge
- ▶ GRU and Conv1D
- ▶ LightGBM
- ▶ Wordbatch FTRL, FM\_FTRL (@anttip)

## 2nd place solution: 0.3889 by Mercaring (Nima & Chahhou)

- ▶ Concatenate brand and category with name
- ▶ Ridge on concatenated name + description: 0.418
- ▶ Sparse NN
- ▶ fastText NN, shared name and description embeddings
- ▶ Sparse NN on a different dataset
- ▶ Double batch size after each epoch

3rd place solution: 0.3905 by @whitebird

- ▶ CNN model: 0.400
- ▶ Wordbatch FM\_FTRL: 0.415
- ▶ A lot of effort on optimization

@sergeif magic feature: any model (inc. Ridge) to  $<0.410$

- ▶ brand  $\times$  name, brand  $\times$  description
- ▶ category  $\times$  name, category  $\times$  description
- ▶ etc ...

brand="Apple", name="iPhone 8s new"  $\Rightarrow$   
{ "Apple\_iPhone", "Apple\_8s", "Apple\_new" }



# Main differences of our approach

- ▶ One model kind, 3 datasets
- ▶ Train 12 models
- ▶ Sparse MLP model
- ▶ Early merge: almost all good ideas created after merging

# Questions?

## First Layer Hidden Size

Hidden size	Score (delta)
128	0.3757 (+0.0024)
256	0.3733 (+0.0000)
384	0.3728 (−0.0005)

## Binarized Features, Classification

Setup	Score (delta)
default	0.3733 (+0.0000)
no binary	0.3740 (+0.0007)
no clf	0.3742 (+0.0009)
no both	0.3748 (+0.0015)